

## Domains:

1. Development with AWS Services  
Serverless, API, SDK, CLI
2. Deployment  
CI/CD, Beanstalk, Serverless
3. Security  
IAM, KMS, each service access policies and encryption methods
4. Monitoring and troubleshooting  
CloudWatch, CloudTrail, X-Ray
5. Refactoring  
AWS Service for best migration

## IAM and EC2:

EC2 is region specific service, while IAM is a global service

IAM: Identity and Access Management

AWS Security: Users, Groups, Roles

Root account should never be used

All AWS services are connected to IAM for security

Policies are written in JSON

IAM controls: Who (authentication) can do What (authorization) in your AWS account.

Authentication(who) with IAM is done with users/groups and roles whereas authorization(what) is done by policies.

User - End user, think about people

Groups - a set of users under one set of permission(policies)

Roles - are used to grant specific permission to specific actors for a set of duration of time. These actors can be authenticated by AWS or some trusted external system.

Can have a password? user->Yes , role->No

Can have an access key? user->Yes , role->No

Can belong to a group? user->Yes , role->No

Can be associated with AWS resources (for example EC2 instances)? user->No , role->yes

AWS supports 3 Role Types for different scenarios:

AWS service roles (for example: EC2, Lambda, Redshift,...)

Cross-Account Access: granting permissions to users from other AWS account, whether you control those accounts or not.

Identity Provider Access: granting permissions to users authenticated by a trusted external system. AWS supports two kinds of identity federation: - Web-based identity such as Facebook, Google- IAM support integration via OpenID Connect and SAML 2.0 identity such as Active Directory, LDAP.

To use/integrate your company credentials, use IAM Federation. It uses SAML standard

#### [IAM best practices:](#)

One IAM user per one person. IAM user should not be assigned to multiple people

One IAM role per application. IAM role should not be assigned to multiple applications

IAM credentials should not be in the code

Don't use root account except for initial setup

IAM policy: service(AWS), actions(functions), resources(specific resource in that service or all the resources in that service for the account), request conditions  
can be assigned to user, group, or role

Use the Principal element in a policy to specify the principal that is allowed or denied access to a resource.

It must not be included in identity-based policies, the principal is implicitly the user that the policy is attached to (for IAM users) or the user who assumes the role (for role access policies)

The principal\_block element is required in resource-based policies (for example, in Amazon S3 bucket policies) and in trust policies for IAM roles.

Resource-based policies are policies that you embed directly in an AWS resource. For example, you can embed policies in an Amazon S3 bucket or an AWS KMS customer master key (CMK).

In resource-based policies, use the Principal element to specify the accounts or users who are allowed to access the resource

In IAM roles, use the Principal element in the role's trust policy to specify who can assume the role. For cross-account access, you must specify the 12-digit identifier of the trusted account.

to attach the policy, use the [attach-user-policy](#) command, and reference the environment variable that holds the policy ARN.

```
$ aws iam attach-user-policy --user-name MyUser --policy-arn $POLICYARN
```

Verify that the policy is attached to the user by running the [list-attached-user-policies](#) command.

```
$ aws iam list-attached-user-policies --user-name MyUser
```

enable your applications to use temporary security credentials provided by AWS by attaching an IAM role to an existing EC2 instance

attach an IAM role to an existing EC2 instance

Call the `associate-iam-instance-profile` command to attach the instance profile,

YourNewRole-Instance-Profile, for the newly created IAM role, YourNewRole, to your EC2 instance, YourInstanceId

```
$aws ec2 associate-iam-instance-profile --instance-id YourInstanceId --iam-instance-profile Name=YourNewRole-Instance-Profile
```

Renting virtual machines - EC2  
Storing data on virtual drives - EBS  
Distributing load across machines - ELB  
Scaling the services using an auto-scaling group - ASG

Create/Launch EC2 instance

- Choose Machine Image (OS)
  - Choose Instance type (CPU, RAM, etc)
  - Configure Instance (VPC/Network, IAM role, File system, User data)
  - Storage (Size, SSD)
  - Tags (key:value pair)
  - Configure Security Group (Firewall around your instance, enable SSH on port 22)
- generate public private key pair, and download it

custom AMI are specific to a region

when trying to ssh into the instance we get permission denied. Permission 0644 for "private key file" is too open.

```
chmod 0400 <file name>
```

```
ssh -i <key file name> user@ip_address
```

connection timeout issue

This is a security group issue. Any timeout (not just for SSH) is related to security groups or a firewall.  
connection refused

This means the instance is reachable, but no SSH utility is running on the instance

permission denied (publickey,gssapi-keyex,gssapi-with-mic)

This means either two things: You are using the wrong security key or not using a security key, or You are using the wrong user

public IP of your EC2 instance will change on instance restart

ec2 connect -> aws service to ssh into your instance from browser

Security Groups: network security in AWS

control how traffic is allowed in and out of EC2 machine

access to ports, authorised ip range, control of inbound and outbound traffic

they are specific to region/vpc combination

all inbound traffic is blocked by default

all outbound traffic is authorised by default

one security group can be connected to another security group - usually for communication between multiple ec2 instances

if you need a fixed public ip for your instance -> Elastic IP. can only have 5 elastic ip in your account but best practice - use random public ip and register dns name to it or use a load balancer

By default, ec2 comes with a private ip (for internal AWS network) and a public IP (www, which we use to ssh)

EC2 user data - run the commands when the machine starts (ex - install updates, required softwares) for example, every new instance we create should have apache server installed on it

aws ec2 instance metadata <- command to learn about themselves

<http://169.254.169.254/latest/meta-data>

internal ip, will only run in the ec2 instance

can retrieve IAM role but not the IAM Policy(to test IAM policy use dry run or policy simulator)

EC2 Instance types:

On demand - short workload, predictable pricing

Reserved (min 1 year) - Reserved instance(long workload), Convertible Reserved(long workload with flexible instance), Scheduled Reserved(every Thursday between a particular time period)

Spot instance - short workload, cheap, less reliable

Dedicated - entire physical server

Elastic Network Interfaces (ENI) -> logical component of VPC that represents virtual network card

ENI can have: primary private ipv4, 1 eip per private ipv4, 1 public ipv4, 1 or more security groups, a mac address

can be created independently and moved from the ec2 instance in case of failover

bound to AZ

## ELB and ASG:

Scalability - system can handle greater loads by adapting

High availability - fault tolerant

Elastic Load Balancer (ELB) -> servers that forward internet traffic to multiple servers downstream expose a single point of access (DNS) to your application

regular health checks of the instances

handles failures of downstream instances

can provide ssl termination for the website

can enforce stickiness with cookies

provides high availability across zones

separates public traffic from private

the health check is done on a port and a route, if the response is not 200 then instance is unhealthy

3 kinds of load balancers in aws:

Classic load balancer (old) - HTTP, HTTPS, TCP

Application load balancer - HTTP, HTTPS, WebSocket

Network load balancer - TCP, UDP, TLS

ELBs can be internal(private) or external(public)

Change the security group of the ec2 instance so that http communication is restricted with the elb, whereas elb's security group can communicate with any user

4xx errors are client induced errors

5xx are application induced errors

load balancer error 503 means at capacity or no registered target

if elb can't connect to app then check security groups

ELB access logs will log all access requests (you can debug per request)

CloudWatch Metrics will give you aggregate stats (ex - connections count)

Application load balancer - level 7

load balancing to multiple applications across machines (target groups)

load balancing to multiple applications on the same machine (containers)

supports websocket

supports redirect

supports routing - based on path in url, based on hostname in url, based on query string, header

port mapping to redirect to a dynamic port in ECS

good for microservices and container based application

Target groups could be -

EC2 instances managed by ASG (HTTP)

EC2 Tasks managed by ECS (HTTP)

Lambda functions - HTTP request is translated into a JSON event

IP addresses to private IPs

we get fixed hostname with ALB

we don't see the client IP directly, IP of client in header X-Forwarded-For, port in X-Forwarded-Port, proto in X-Forwarded-Proto

Network Load Balancer - level 4

forward tcp and udp traffic to your instances

handles millions of requests per second, less latency 100ms (vs 400ms of alb)

has one static ip per AZ, supports assigning eip, exposes ip instead of dns

used for extreme performance, tcp or udp traffic

ELB - stickiness

same client redirected to same instance behind a load balancer

works for CLB and ALB

the cookie used for stickiness has an expiration date you control

use case: make sure the user doesn't lose his session data

enabling it may bring imbalance

when cross zone load balancing used, each lb distributes load across all registered instances in all AZ

SSL/TLS - allows traffic between client and load balancer to be encrypted in transit (in-flight encryption)

Secure Sockets Layer (SSL) to encrypt connections

Transport Layer Security (TLS), newer version

load balancer does ssl termination

SNI - server name indication, handle multiple certificates with your load balancer

ELB - connection draining (CLB), Deregistration delay (target group, alb and nlb)

time to complete "in-flight requests" while the instance is de-registering or unhealthy

ASG - Auto scaling group

scale out to match increased load

scale in to match decreased load

automatically register new instances to load balancer

ASG have

launch configuration: AMI + instance type, EC2 user data, EBS Volume, Security groups, SSH key pair

min size, max size, initial capacity

network + subnet info

load balancer/target group info

scaling policies

possible to scale ASG based on CloudWatch alarms

auto scaling rules based on - target avg cpu usage, no. of requests per instance, avg network in, avg network out

ASG - scaling policies

Target tracking scaling - avg. ASG CPU at 50%

Simple/Step scaling - CloudWatch alarm cpu > 70%, then add 2 units, CloudWatch alarm cpu < 30%, then remove 1 units

Scheduled Actions - anticipate based on patterns

Scaling cooldown - asg doesn't take any action before previous scaling activity takes effect

## EBS and EFS:

Elastic Block Storage Volume - network drive attached to instances  
EC2 loses its main drive when terminated, store data somewhere else

latency, can be detached from one instance and attached to another instance  
locked to an AZ, can be moved across AZs using snapshot  
has a provisioned capacity, size in GB and iops

4 types:

GP2 (SSD): General purpose volume

IO1 (SSD): High performance, low latency - large database workloads

ST1 (HDD): Low cost, throughput intensive workloads - log processing, kafka

SCI (HDD): Lowest cost, less frequently accessed

gp2 and io1 can be used as boot volumes

iops: io operation per second

newly created ebs volume needs to be mounted to the ec2 instance, after formatting the volume with  
any file system

billed for provisioned

instance store (ephemeral) is physically attached to the machine, ebs is a network drive

Elastic File System - can be mounted on many EC2 instances in multi-AZ

content management, web serving, data sharing

only for linux instances, posix

billed for usage

instance store locked to ec2 instance

EBS locked to an AZ

EFS locked to multi-AZs (Region?)

to migrate EBS volume, take snapshot and then restore it in another AZ

## RDS, Aurora and ElastiCache:

Relational Database Service: managed db service for dbs using sql  
automated backups, transaction logs, snapshots

read replica: up to 5 read replicas possible

could be with in AZ, Cross AZ, or cross region

reads are eventually consistent

can be promoted to their own db

network cost when data goes from one AZ to another

multi AZ (disaster recovery): sync replication

one dns name, automatic app failover(AZ, network, db, instance) to standby

the read replicas can be setup as multi AZ for disaster recovery

at rest encryption, if master is not encrypted than read replicas can't be encrypted  
in flight encryption, ssl. enforce ssl from database

deployed in private subnet, security groups

iam policies to manage RDS

iam based authentication to log into rds, auth token

Aurora - proprietary aws technology

storage automatically grows in increments of 10 GB, upto 64 TB

15 read replicas

6 copies of your data across 3 AZs

aurora serverless - no capacity planning needed

good for infrequent, intermittent or unpredictable workloads

db instantiation and auto scaling based on usage

global aurora:

Aurora Cross region read replicas:

- useful for disaster recovery

- simple to put in place

Aurora global database:

- 1 primary region

- 5 read only regions

ElastiCache: Managed redis or memcached

in-memory db with high performance and low latency

write scaling using sharding, read scaling using read replicas

multi AZ with failover

can be used to store user session data

Redis: multi AZ with auto-failover

read replicas to scale reads and have high availability

data durability using AOF persistence(logs every write operation received by server)



backup and restore features

Memcached: multi-node for partitioning of data (sharding)

non persistent

no backup, restore

multi threaded architecture

Caching policy:

Lazy loading/Cache-aside/Lazy population : update the cache using missed data on cache miss

Write through : add or update cache when database is updated

Cache evictions and TTL: explicit delete or mem is full and cache is not used (lru) or item TTL expires

## Route53:

managed DNS

common types of records:

A: hostname to ipv4

AAAA: hostname to ipv6

CNAME: hostname to hostname

Alias: hostname to aws resource

public domain names and private domain names

load balancing (DNS - client load balancing)

Health checks (limited func.)

routing policies

pay per month per hosted zone

TTL

if root domain than use alias, for non root domain use alias or cname

simple routing policy: redirect to a single resource, in case of multiple returns client will randomly pick one

weighted routing: control the % of requests that go to specific endpoint

latency routing: least latency from user to designated AWS region

failover routing: use secondary if health check on primary fails

geolocation: route based on user location

multi value routing: return healthy records

## VPC:

virtual private cloud - private network to deploy your resources

Vpcs are region specific, 5 vpc per region

subnets - partition network inside VPC, locked to AZ, 200 subnets per vpc

Region comes with default vpc

public subnet accessible from internet, private not accessible from internet

to define access to internet, and between subnets, use route tables

Each subnet in vpc must be associated with a route table

route table contains a set of rules (routes) that determine where network traffic is directed.

Each subnet in your VPC must be associated with a route table which controls the routing for subnet.

A subnet can only be associated with one route table at a time, but you can associate multiple subnets with the same route table.

internet gateways helps our vpc instances connect with the internet

public subnets have a route to internet gateway

Provides a target in the vpc route table for internet-routable traffic

Perform network address translation (NAT) for instances that have been assigned ipv4

NAT gateways (AWS managed) and NAT instances (self-managed)

allow your instances in your Private Subnets to access internet while remaining private

Network ACL - firewall which controls traffic from and to subnet

allow and deny rules, ip address at subnet level

default everything allowed in and out

Security Group - firewall that controls traffic to and from an ENI/an EC2 instance

allow rules, ip addresses or other security groups

VPC flow logs: capture info about traffic going into your interfaces, VPC, Subnet, ENI (elastic network interface) flow logs

VPC peering: connect two VPC, privately using AWS network, must not have overlapping cidr block (IP address range), not transitive

VPC endpoints: connect to AWS services using private network instead of www, vpc endpoint gateway (s3, dynamoDB), vpc endpoint interface (other services)

Site to site vpn: connect on prem VPN to AWS over internet

Direct connect: physical connection

## S3:

objects - files, buckets - directories  
buckets must have globally unique name  
S3 global service but buckets are region level

objects(files) have a key (Full path), prefix + object name  
max object size is 5 TB, to upload more than 5 GB use multi-part upload

metadata  
tags

Versioning: enabled at bucket level, version your files in S3  
same key overwrite will increase the version  
versions of a file: queue of versions but emulates a stack with head being the latest/current version

pre signed url: can be generated using CLI(download) or SDK(upload)  
valid for default 3600 seconds  
generate URL dynamically to download file instead of giving user the access to the bucket or give temporary access  
cli commands:

```
aws configure set default.s3.signature_version s3v4  
aws s3 presign s3://<object-path> --region <location>
```

Encryption: encrypt objects, 4 methods, SSE (server side encryption)  
SSE-S3 : encrypt s3 objects using keys handled and managed by AWS, set header  
"x-amz-server-side-encryption":"AES256"  
SSE-KMS : use AWS-KMS to manage encryption keys, user control and audit trail,  
"x-amz-server-side-encryption":"aws:kms"  
SSE-C : manage your own encryption keys, https, encryption key in every http header, aws will discard the key, only thru cli  
Client side encryption, encrypt data before sending, Amazon S3 encryption client  
encryption in transit - ssl/tls, amazon s3 exposes https endpoint (mandatory for SSE-C), http endpoint not encrypted

S3 put api headers:

Content-MD5 - This header can be used as a message integrity check to verify that the data is the same data that was originally sent.

X-amz-server-side-encryption - The server-side encryption algorithm used when storing this object in Amazon S3 (for example, AES256, aws:kms).

X-amz-server-side-encryption-aws-kms-key-id - If the value of x-amz-server-side-encryption is aws:kms, this header specifies the ID of the symmetric customer managed AWS KMS CMK that will be

used for the object. If you specify `x-amz-server-side-encryption:aws:kms`, but do not provide `x-amz-server-side-encryption-aws-kms-key-id`, Amazon S3 uses the AWS managed CMK in AWS to protect the data.

`X-amz-server-side-encryption-customer-key` - Specifies the customer-provided encryption key for Amazon S3 to use in encrypting data. This value is used to store the object and then it is discarded; Amazon S3 does not store the encryption key. The key must be appropriate for use with the algorithm specified in the `x-amz-server-side-encryption-customer-algorithm` header.

`x-amz-server-side-encryption-customer-key-MD5` - Specifies the 128-bit MD5 digest of the encryption key according to RFC 1321. Amazon S3 uses this header for a message integrity check to ensure that the encryption key was transmitted without error.

`X-amz-storage-class` - If you don't specify, S3 Standard is the default storage class. Amazon S3 supports other storage classes.

User based S3 security: IAM policies - which api calls should be allowed from which user

Resource based S3 security: bucket policies, Object Access Control List (ACL), Bucket Access Control List (ACL)

Bucket ACLs allow you to control access at a bucket level, while Object ACLs allow you to control access at the object level.

json based bucket policies

Resources: buckets and objects

Actions: set of API calls to be Allow or Deny

Effect: Allow or Deny

Principal: the account/user to apply the policy to

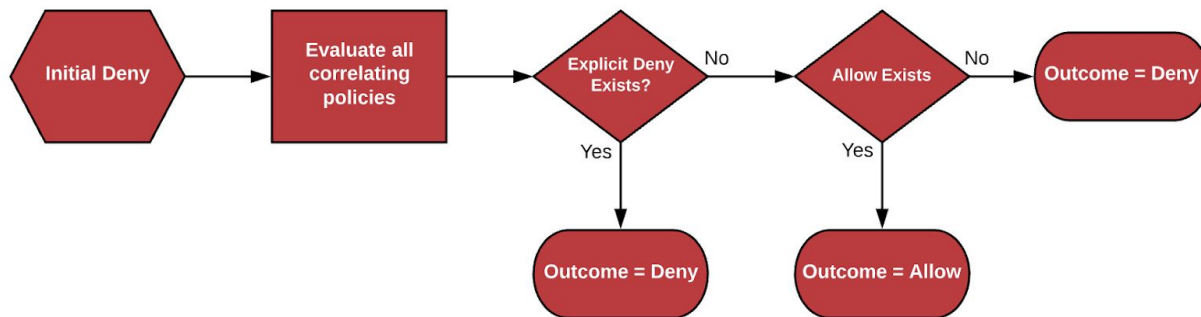
Bucket policies allow users to easily grant cross-account access without having to create roles using the "Principal" IAM element. To see this in action, let's assume we want to allow `s3:GetObject` on `foobucket` to the root account `6161616161` as well as the user `jason` under that account. This could be implemented with the bucket policy below:

```
{
  "Version": "2012-10-17",
  "Statement": [{
    "Sid": "Allow",
    "Effect": "Allow",
    "Principal": {
      "AWS": [
        "arn:aws:iam::616161616161:root",
        "arn:aws:iam::616161616161:user/jason"
      ]
    },
    "Action": "s3:GetObject",
    "Resource": "arn:aws:s3:::foobucket/*"
  }]
}
```

the user based policy below allows whatever entity that has the policy attached to perform any action on foobucket:

```
{
  "Version": "2012-10-17",
  "Statement": [{
    "Effect": "Allow",
    "Action": "s3:*",
    "Resource": ["arn:aws:s3:::foobucket",
      "arn:aws:s3:::foobucket/*"]
  }]
}
```

access will be determined by combination of all policies:



Use IAM policies if:

- You need to control access to AWS services other than S3. IAM policies will be easier to manage since you can centrally manage all of your permissions in IAM, instead of spreading them between IAM and S3.

- You have numerous S3 buckets each with different permissions requirements. IAM policies will be easier to manage since you don't have to define a large number of S3 bucket policies and can instead rely on fewer, more detailed IAM policies.

- You prefer to keep access control policies in the IAM environment.

Use S3 bucket policies if:

- You want a simple way to grant [cross-account access](#) to your S3 environment, without using [IAM roles](#).

- Your IAM policies bump up against the size limit (up to 2 kb for users, 5 kb for groups, and 10 kb for roles). S3 supports bucket policies of up to 20 kb.

- You prefer to keep access control policies in the S3 environment.

put-bucket-policy applies an Amazon S3 bucket policy to an Amazon S3 bucket.

```
aws s3api put-bucket-policy --bucket MyBucket --policy file://policy.json
```

If you are using an identity other than the root user of the AWS account that owns the bucket, the calling identity must have the PutBucketPolicy permissions on the specified bucket and belong to the bucket owner's account in order to use this operation.

If you don't have PutBucketPolicy permissions, Amazon S3 returns a 403 Access Denied error. If you have the correct permissions, but you're not using an identity that belongs to the bucket owner's account, Amazon S3 returns a 405 Method Not Allowed error.

grant public access to the bucket, force objects to be encrypted at upload, grant access to another account

bucket setting to block public access, when you know bucket should not be public

bucket policies are evaluated before default encryption

S3 network security - vpc endpoints (for instances in vpc without internet)

Logging and audit - S3 access logs can be stored in other s3 bucket, API calls can be logged in AWS CloudTrail

User security - MFA Delete, pre signed urls valid for a limited time

S3 websites - host static websites, accessible from internet, 403 error if bucket policy denies access

path = bucket/prefix/number1/file01, bucket + prefix + file

Cross-Origin Resource Sharing ([CORS](#)) is a mechanism that uses additional [HTTP](#) headers to tell browsers to give a web application running at one [origin](#), access to selected resources from a different origin. A web application executes a cross-origin HTTP request when it requests a resource that has a different origin (domain, protocol, or port) from its own.

Cross Origin Resource Sharing (CORS):

- origin: scheme (protocol), host (domain), port

- get resources from different origin

- browser based mechanism to allow requests to other origins while visiting the main origin

- request won't be fulfilled unless the other origin allows, using CORS header

Access-Control-Allow-Origin, Access-Control-Allow-Methods

- preflight request/response

- for client to do cross-origin request on S3, we need to enable the headers

- can allow specific origin or \*

- s3 bucket -> permissions -> CORS configuration -> paste CORSRule

S3 consistency: eventually consistent

read after write consistency for PUTS of NEW objects, except if you had done get before put

eventual consistency for DELETE and PUT of EXISTING objects

-> S3 and Athena

MFA-Delete, first enable versioning in S3 bucket, then permanently deleting object version and suspending versioning on the bucket will require MFA

only root account/bucket owner can enable/disable MFA-Delete, using CLI

S3 access logs: all requests logged into another s3, can be analysed using athena

S3 replication: cross region replication(crr), Same region replication(srr)

versioning should be enabled in source and destination, buckets can be in different account,

asynchronous replication, proper IAM permissions

only new objects after enabling will be replicated, delete not replicated, no chaining of replication

CRR: compliance, lower latency access, replication across accounts

SRR: log aggregation, live replication b/w prod and test

S3 storage classes of object:

s3 standard - general purpose : high durability, across multi AZ, for big data analytics, mobile and gaming applications, content distribution

s3 standard infrequent access (IA): for less frequent but rapid accessed, disaster recovery, backups

s3 one zone infrequent access: same as IA but in single zone, 20% less cost, secondary data or data you can recreate

s3 intelligent tiering: low latency and high throughput same as standard, automatically moves objects between two tiers

glacier: low cost for archiving/backup, longer terms (10 years), 3 retrieval options (expedited(1-5 minutes), standard(3 to 5 hours), bulk(5 to 12 hours)), min 90 days of storage

glacier deep archive: 2 retrieval options (standard(12 hours), bulk(48 hours)), min 180 days of storage

S3 - object lifecycle: move objects between storage classes

infrequently accessed object to standard ia

archive objects not needed in the real time to glacier

automated using lifecycle configuration

transition actions, expiration actions

s3 performance might get restricted by KMS if you use sse-kms

to improve s3 performance for upload: multi-part upload for large files, s3 transfer acceleration (transfer to s3 edge location)

s3 byte range fetches - parallelize get

server side filtering - select query in s3 and glacier

s3 event notification - trigger based on s3 event, can filtered based on object name as well before triggering

target of event notification - sns, sqs, lambda

version is required whenever you use policy variables

```
{
  "Sid": "AllowAllS3ActionsInUserFolder",
  "Action":["s3:*"],
  "Effect":"Allow",
  "Resource":["arn:aws:s3:::my-company/home/${aws:username}/*"]
}
```

Whenever a user makes a request to AWS, the variable is replaced by the "friendly" user name of whomever made the request.

aws athena - serverless analytics, directly against s3 files

s3 object lock and glacier vault lock - write once read many, block object version deletion/edit for a time period

## AWS CLI, SDK:

perform interactions with aws without online console



interact with aws proprietary service

aws tasks against aws can be done using

- AWS CLI on local computer

- AWS CLI on EC2 instance

- AWS SDK on local computer

- AWS SDK on EC2 instance

- AWS instance metadata service for EC2

during cli installation aws executable's path is generally added to the environment variables

cli configure: aws access key and Secret key (don't share)

get user's access keys, IAM -> Users -> user profile -> Security Credentials -> Create access key -> download

aws configure <- to configure user access, access key id, secret access key, default region; will create config and credentials files in ~/.aws

aws cli on EC2 using IAM roles: assign IAM roles with policy authorizing what the EC2 instance can do commands on cli:

aws s3 ls

aws s3 <command\_name> help

aws cli dry run (simulate call, check permission) flag --dry-run. On success, the request would have succeeded but the DryRun flag is set.

error message can be decoded using STS command line

aws sts decode-authorization-message --encoded-message <err\_msg>

aws ec2 instance metadata - learn about themselves

<http://169.254.169.254/latest/meta-data>

internal ip, will only run in the ec2 instance

can retrieve IAM role but not the IAM Policy(to test IAM policy use dry run or policy simulator)

to use MFA with cli, create a temporary session. run STS GetSessionToken API call

aws sts get-session-token --serial-number arn-of-the-mfa-device --tokencode code-from-token

--duration-seconds 3600

configure the profile and add aws\_session\_token in the credentials file

perform actions on AWS from your applications

AWS SDK: java, .net, node, python, go, etc.

AWS CLI uses python SDK

when default region is not specified, us-east-1 is chosen

aws limits:

API rate limits

DescribeInstances API for EC2 has a limit of 100 calls per seconds  
GetObject on S3 has a limit of 5500 GET per second per prefix  
KMS API has rate limit  
For Intermittent Errors: implement Exponential Backoff  
For Consistent Errors: request an API throttling limit increase

#### Service Quotas (Service Limits)

Running On-Demand Standard Instances: 1152 vCPU  
to request a service limit increase open a ticket  
request a service quota increase using the Service Quotas API

credentials provider chain(priority): command line options, env variables, cli credentials file, cli configuration file, container credentials, instance profile credentials  
for sdk: env variables, system properties, default credential profiles file, ecs container credentials, instance profile credentials

all of our aws http api requests should be signed(AWS credentials). CLI and SDK do that for us  
otherwise use SigV4

SigV4: HTTP header, or Query String Option(pre-signed URLs)

## AWS CloudFront:

Content Delivery Network (CDN), improves read performance, content is cached at the edge  
ddos protection, integration with shield, aws web application firewall  
can expose external https and talk to internal https backend

cloudfront origins - s3 bucket, custom origin(http) - alb, ec2, s3 website, any http backend  
cloudfront geo restriction - whitelist, blacklist countries

cloudfront uses global edge network and cached for a ttl, good for static content  
vs s3 crr setup for each region, files updated real-time, read only, good for dynamic content available to few regions at low latency

origin access identity - users access s3 content only using cloudfront urls

cache based on header, session cookies, query string parameters at edge locations  
expire based on ttl  
cache can be invalidated using api call  
maximize cache hits by separating static(s3) and dynamic(alb, based on header and cookie)  
distributions

viewer protocol policy: between client and edge location - redirect http to https, https only  
origin protocol policy: between origin and edge location - https only, match viewer

s3 bucket websites don't support https

signed url/cookie: attach a policy (url expiration, IP ranges to access data from, trusted signers)

signed url: one signed url per file, signed cookie: one signed cookie for multiple files

CloudFront Signed URL:

Allow access to a path, no matter the origin. Account wide key-pair, only the root can manage it.

Can filter by IP, path, date, expiration. Can leverage caching features

S3 Pre-Signed URL:

Issue a request as the person who pre-signed the URL. Uses the IAM key of the signing IAM principal. Limited lifetime

## **ECS, ECR and Fargate:**

Docker containers

ECS: cluster, services, tasks, tasks definition

dockerfile - docker configuration

docker image - built from dockerfile

docker hub - stores docker images

docker container - running instance of docker image

to manage containers, we have container management platform: ECS, Fargate, EKS

ecs clusters are logical grouping of EC2 instances, EC2 runs a special AMI made for ECS

EC2 instances run ECS agent(docker container), these ecs agent registers the instance to ecs cluster

ecs task definition - metadata in json format, tell ecs how to run a docker container

has image name, port binding for container and host, memory and cpu required, env variables, networking info, etc

specify which containers are included in your task and how they interact with each other

ecs services define how many tasks should run and how, ensure no of tasks desired is running across ec2 instances

can be linked to load balancers

dynamic port forwarding in load balancer, use it with ecs

ECR: private docker image repository, accessed through IAM,

cli command: `aws ecr get login password` | `docker login` previous command's o/p

`docker push` and `docker pull` command

Fargate: in ecs cluster we have to create ec2 instances on our own

in fargate we just create task definitions

ECS agent needs a EC2 role to attached to it to function, ecs tasks need role attached based on the services they interact with

task placement strategy/task placement constraints to determine where to place or remove tasks from  
task placement strategy: binpack(least available amount of cpu or memory), random, spread (evenly based on parameter(AZ, instance) )

task placement constraint: distinct instance, member of(if satisfy the expression than place there)

ECS service auto scaling based on cloudwatch, target tracking, step scaling, scheduled scaling  
ecs service auto scaling does not mean ec2 auto scaling  
for that ecs cluster capacity provider

To properly instrument your applications in Amazon ECS, you have to create a Docker image that runs the X-Ray daemon, upload it to a Docker image repository, and then deploy it to your Amazon ECS cluster.

You can use port mappings and network mode settings in your task definition file to allow your application to communicate with the daemon container.

## **AWS Elastic Beanstalk:**

deploying application on aws  
instance configuration and os is handled by beanstalk  
deployment strategy is configurable but performed by EB

three architecture models:  
single instance deployment: good for dev  
LB + ASG: for production or pre-production web app  
ASG only: for non web app in production (workers) (process msgs of a queue)

EB has three components:  
application  
application version (each deployment)  
env name (free naming: dev, test, prod)

deploy application versions to environments and can promote to next environment  
rollback, control over lifecycle of environments

elastic beanstalk deployment modes:  
all at once  
rolling  
rolling with additional batches

immutable  
swap url, blue/green deployment

beanstalk lifecycle policy based on time or space.  
elastic beanstalk uses cloudformation in the back  
you can clone eb environment, ex. to deploy test version  
beanstalk migration: to change lb or decouple db

### Beanstalk with HTTPS

- Load the SSL certificate onto the Load Balancer
- Can be done from the Console (EB console, load balancer configuration)
- Can be done from the code: `.ebextensions/securelistener-alb.config`
- SSL Certificate can be provisioned using ACM (AWS Certificate Manager) or CLI
- Must configure a security group rule to allow incoming port 443 (HTTPS port)

### Beanstalk redirect HTTP to HTTPS

- Configure your instances to redirect HTTP to HTTPS:  
<https://github.com/awsdocs/elastic-beanstalk-samples/tree/master/configuration-files/aws-provided/security-configuration/https-redirect>
- OR configure the Application Load Balancer (ALB only) with a rule
- Make sure health checks are not redirected (so they keep giving 200 OK)

For custom configuration files which are not readily available in Elastic Beanstalk, create YAML- or JSON-formatted documents with a `.config` file extension that you place in a folder named `.ebextensions` and deploy in your application source bundle.

`Dockerrun.aws.json` is used in multicontainer Docker environments that are hosted in Elastic Beanstalk  
`env.yaml` is used to configure the environment name, solution stack, and environment links to use when creating your environment in Elastic Beanstalk.

`cron.yaml` is primarily used to define periodic tasks that add jobs to your worker environment's queue automatically at a regular interval

`Appspec.yml` (application specification) is used to manage each application deployment as a series of lifecycle event hooks in CodeDeploy and **NOT in Elastic Beanstalk**.

## CICD - CodeCommit, CodePipeline, CodeBuild, CodeDeploy:

push our code in a repository and have it deployed on AWS

codecommit: storing our code

codepipeline: automating our pipeline from code to elastic beanstalk

codebuild: building and testing our code

codedeploy: deploying the code to ec2

codecommit authorization in Git: IAM Policies manage user / roles rights to repositories

Cross Account access: Use IAM Role in your AWS Account and use AWS STS (with AssumeRole API)  
trigger notifications using AWS SNS or AWS Lambda or AWS CloudWatch event rules

We can migrate a Git repository to a CodeCommit repository in a number of ways: by cloning it, mirroring it, migrating all or just some of the branches, and so on. You can also migrate local, unversioned content on your computer to CodeCommit.

codepipeline -> visual workflow

sequential or parallel actions, build/ test / deploy / loadtest, stages can have multiple action groups

each pipeline stage can create artifacts, which are stored in S3 and passed to next stage

codepipeline state change -> AWS CloudWatch Events -> SNS notifications

aws cloudtrail can be used to audit aws api calls

if pipeline can't perform an action, IAM service role

codebuild - build instructions can be defined in buildspec.yml(phases: install, pre build, build, post build)  
output logs to S3 and CloudWatch Logs, detect failed builds and trigger notifications using CloudWatch alarms

cloudwatch events and AWS lambda, sns notification

build can be defined in codepipeline instead of codebuild as well

buildspec.yml is the file name for the file that defines the build instructions for AWS CodeBuild.

codedeploy - appspec.yml

file section - how to source them (s3/github)

hooks - instructions to deploy new version

- applicationStop

- downloadBundle

- beforeInstall

- afterInstall

- applicationStart

- validateService

- beforeAllowTraffic

- allowTraffic

- afterAllowTraffic

config: one at a time, half at a time, all at once, custom

The name of the AppSpec file for an EC2/On-Premises deployment must be appspec.yml. The name of the AppSpec file for an Amazon ECS or AWS Lambda deployment must be appspec.yaml.

EC2/On-Premises: in-place or blue/green deployment

AWS Lambda: canary, linear, or all-at-once; traffic is shifted to the updated Lambda function versions during a deployment

Amazon ECS: blue/green deployment, traffic shifting

codestar : dashboard/wrapper around all the application management services

## CloudFormation:

infrastructure as code, provisioning mechanism  
template(json or yaml) uploaded in s3, can't be updated, upload new version  
[cloudformation template](#): resource(services), parameters(!Ref function to refer parameters or resources), mapping(switch case, !FindInMap),  
output(optional, can import into other stacks, link stacks, !ImportValue),  
conditionals(environment, region, etc. ), metadata  
create, update and delete stacks using this  
some unsupported aws services can be created using AWS Lambda custom resources  
changeSets, nested stacks, stackSets (across multiple accounts and regions)

Change sets allow you to preview how proposed changes to a stack might impact running resources  
StackSets allows us to provision a common set of AWS resources across multiple accounts and regions with a single cloudformation template

## CloudWatch, X-Ray, CloudTrail:

CloudWatch:

metrics: collect and track key metrics  
logs: collect, monitor, analyze and store log files  
events: send notifications when certain events happen in AWS  
alarms: react in real time to metrics/events

X-Ray:

troubleshooting app and performance  
distributed tracing of microservices

CloudTrail:

internal monitoring of API calls being made  
audit changes to AWS resources by your users

metric, variable being monitored, cpuUtilization, networkIn  
they belong to a namespace(categories), dimension(instance id, environment, etc.) is an attribute of the metric, upto 10 dimensions per metric  
metrics have timestamps  
detailed monitoring for extra cost to increase frequency (ec2 every 1 minute), ec2 ram usage is not pushed as a metric (make custom push function)  
higher metric resolution (storageResolution API parameter), send metric to cloudWatch (PutMetricData api call), use exponential backoffs in case of throttle errors

alarm trigger notifications for any metric, go to Auto scaling, EC2 actions, sns notif.

they have various options (sampling, %, max, min)  
alarm states (ok, insufficient\_data, alarm), period (length of time in seconds to evaluate the metric)

applications can send logs to cloudwatch using sdk  
collects log from - beanstalk, ec2, lambda, vpc flow logs, api gateway, cloudtrail, log agents, route53  
can be analysed, batch export to s3, stream to elasticSearch  
can use filter expressions  
log group(application), log stream(instances within application), log expiration policies  
use correct IAM policies to send logs to cloudwatch, encryption at group level using KMS

cloudwatch logs agent, cloudwatch unified agent  
metric filter

eventbridge: partner event bus (receive from saas), custom event bus (for your application), rules to process events, schema registry(analyze, structure)

X-ray (debugging), visual analysis, understand dependencies, identify bottlenecks, time sla, find errors and exceptions  
tracing (follow a request), each component adds its own trace, segments, sub segments, sampling  
IAM authorisation, KMS encryption

code imports the aws x-ray sdk (application sdk will capture calls to aws services, http/https requests, db calls, queue calls)

install the x-ray daemon or enable x-ray aws integration (should have IAM permissions, lambda run them for us, low level UDP packet interceptor on port 2000)

instrument application code

trace (segments collected together, end-to-end request, X-Amzn-Trace-Id), segment(each service/app will send them), subsegments (for granular details in segment)

annotations: key value pairs used to index traces and use with filters

Trace sampling rate: first request each second, five percent additional requests

subsegment can contain additional details about a call to an AWS service, an external HTTP API, or an SQL database. You can define arbitrary subsegments to instrument specific functions or lines of code in your application.

For services that don't send their own segments like DynamoDB, X-Ray can use subsegments to generate inferred segments and downstream nodes on the service map

Segment documents can be up to 64 kB in size

AWS X-Ray daemon is a software application that listens for traffic on UDP port 2000, gathers raw segment data, and relays it to the AWS X-Ray API

X-Ray takes the client IP from the X-Forwarded-For header in the request instead of from the source IP in the IP packet

To configure X-ray SDK use environment variables daemon address(AWS\_XRAY\_DAEMON\_ADDRESS – Set the host and port of the X-Ray daemon listener. By default 127.0.0.1:2000 for both trace data (UDP) and sampling (TCP)), context missing(AWS\_XRAY\_CONTEXT\_MISSING – Set to



LOG\_ERROR to avoid throwing exceptions when your instrumented code attempts to record data when no segment is open.), trace id(AWS\_XRAY\_TRACING\_NAME, used be segment)

#### X-Ray Write APIs:

PutTraceSegments: Uploads segment documents to AWS X-Ray

PutTelemetryRecords: Used by the AWS X-Ray daemon to upload telemetry. Ex:

SegmentsReceivedCount, SegmentsRejectedCounts, BackendConnectionErrors

GetSamplingRules: Retrieve all sampling rules

#### X-Ray Read APIs:

GetServiceGraph: main graph

BatchGetTraces: Retrieves a list of traces specified by ID. Each trace is a collection of segment documents that originates from a single request.

GetTraceSummaries: Retrieves IDs and annotations for traces available for a specified time frame using an optional filter. To get the full traces, pass the trace IDs to BatchGetTraces.

GetTraceGraph: Retrieves a service graph for one or more specific trace IDs.

cloudtrail: enabled by default, history of event/api calls made by the aws account  
something deleted look into cloudtrail

## SQS, SNS, Kinesis:

asynchronous/event based communication, app -> queue -> app, applications decoupled

sqs: queue model, sns: pub/sub model, kinesis: real-time streaming

producers send messages -> sqs queue -> consumers poll messages

4 days message retention

can have duplicate messages(at least once delivery), out of order messages(best offer ordering)

256kb per message

delay queue - by default or using message parameter(DelaySeconds)

poll sqs messages by consumer (receive upto 10 at a time)

have to process the message within visibility timeout, then delete(DeleteMessage) from queue using msg id

if consumer needs more time to process a message than tell the queue to increase the timeout (ChangeMessageVisibility api call)

dead letter queue - if the consumer fails to process a message multiple times(threshold, redrive policy), send the message to DLQ (dead letter queue that you created) from main sqs

sqs - long polling, if consumer doesn't find any messages in the queue it will wait for the wait time specified. decreases api calls on empty queues

enable by default or api call WaitTimeSeconds

sqs fifo queue - name must end in .fifo  
lower throughput, no per msg delay, sent exactly once  
MessageDeduplicationId - to avoid duplicates  
sequencing - MessageGroupId, that many consumers

sqs extended client (java library)  
sending message larger than 256kb -> producer sends it to s3 and metadata to queue -> consumer receives metadata and extracts file from s3

SNS - one message to many receivers, fan out

Kinesis - big data streaming tool  
kinesis streams - low latency streaming ingest at scale  
streams are divided into shards/partitions  
data can be reprocess, multiple apps can consume the same stream  
data is immutable  
number shards = number of max kcl possible  
kinesis analytics - perform real time analytics on streams using sql  
kinesis firehose - load streams into s3, redshift, elasticsearch

## AWS Lambda:

no need to manage servers, just deploy the code  
lambda - virtual functions, auto scaled and on-demand, short executions  
lambda + api gateway : rest api  
lambda + kinesis : data transformations on the fly  
lambda + dynamodb : database event related processing  
lambda + s3 : s3 event can invoke lambda function  
lambda@edge + cloudfront :  
lambda + cloudwatch events EventBridge : invoke lambda based on our architecture state change  
lambda + cloudwatch :  
lambda + sns : react to notification  
lambda + sqs : process sqs queue messages  
lambda + cognito : react when user logs in

example: create a thumbnail on the fly - new image in s3 -> trigger lambda which creates a thumbnail -> stores in another s3 and inserts data to dynamodb  
example: serverless cron job - cloudwatch events EventBridge triggers every one hour -> lambda function performs a task

package code and deployment dependencies, upload it to create lambda function, lambda stores code in s3 and encrypts it at rest

a layer is a ZIP archive that contains libraries, a custom runtime, or other dependencies. use libraries in your function without including them in the deployment package.

lambda function invocation can be synchronous or asynchronous (on-demand invocation)

synchronous - wait for the result, which is returned right away, when you use cli, sdk, api gateway, application load balancer

error handling on client side

user invoked(alb, api gateway, cloudfront, s3 batch), service(cognito, step function), other(lex, alexa, kinesis firehose) are synchronous

lambda function asynchronous invocation by s3 event, sns, cloudwatch events/eventbridge, codecommit trigger, codepipeline, etc

from cli set flag --invocation-type Event

lambda reads from an event queue, 3 tries on errors after that it should be pushed to dead letter queue

event source mapping, invoke your lambda function when that event occurs

necessary when the records need to be polled (not pushed events), kinesis data streams, sqs, dynamodb streams

lambda function itself synchronous invocation

streams - iterator for each shard, batch records, on error batch is reprocessed, destination(send the result of asynchronous invocation, upgrade compared to dlq)

queue - long polling, batch size, dlq

lambda event mapper, needs permission to read from sqs

by default, AWS executes Lambda function code securely within its own VPC

lambda with ENI

Deploying a Lambda function in a public subnet does not give it internet access or a public IP

Deploying a Lambda function in a private subnet gives it internet access if you have a NAT Gateway / Instance

You can use VPC endpoints to privately access AWS services without a NAT

Triggering lambda function:

Synchronous: RequestResponse

user invoked(alb, api gateway, cloudfront, s3 batch), service(cognito, step function), other(lex, alexa, kinesis firehose)

Event source mapping:

Poll trigger: dynamodb, kinesis, sqs

configuration of the event source mapping is on the Lambda side

Asynchronous: Event

Push trigger: s3, sns, cloudwatch events/eventbridge, codecommit trigger, codepipeline

configuration is made on the source (S3/SNS)

the number of concurrent executions for poll-based event sources is different from push-based event sources

iam role vs resource based policy

to expose lambda function as HTTP(S) endpoint, use application load balancer (or api gateway)

lambda function should be registered in target group

json to http and http to json conversion in alb, and multi header value in target group setting

Lambda@Edge - deploy lambda functions along side cloudfront cdn

run lambda functions to customize content that CloudFront delivers, execute functions in AWS locations closer to the viewer.

functions run in response to CloudFront events

use lambda functions to change CloudFront requests and responses at the following points:

after CloudFront receives a request from a viewer (viewer request)

before CloudFront forwards the request to the origin (origin request)

after CloudFront receives the response from the origin (origin response)

before CloudFront forwards the response to the viewer (viewer response)

lambda environment variables, encrypt using kms

lambda automatically monitors functions on our behalf, reporting metrics through Amazon CloudWatch  
CloudWatch metrics include total invocation requests, duration, concurrent executions and error rates, throttles, async delivery failures

throttles, Dead Letter Queues errors and Iterator age for stream-based invocations are also monitored

CloudWatch does not monitor lambda's reserved concurrent executions but can view it through the lambda console

active tracing, lambda tracing with x-ray

lambda - upto 1000 concurrent executions, can be limited using "reserved concurrency", invocation over the limit will trigger a throttle

throttle in synchronous invocation - return `throttledError` 429

throttle in asynchronous invocation - retry and then go to dlq

first request served by new instances has higher latency than rest - cold start

provisioned concurrency - allocated before the function is invoked, cold start never happens

lambda function dependencies - install the package along your code and zip it together

if zip less than 50 MB then upload to lambda else upload to s3 and reference it

custom runtime or reusable dependencies using lambda layers

create connection to other services outside of function (context), so that it can be reused during invocations

lambda execution limits per regions - max execution time 15 min, memory 128MB to 3TB (64 MB increments),  
env variables 4KB, disk capacity in function container (/tmp) 512MB <- local directory, concurrency executions 1000  
lambda deployments limits - size of compressed zip file 50MB, size of uncompressed deployment (code+dependencies) 250MB,  
can use /tmp directory to load other files at startup, env variables 4KB

## DynamoDB:

highly available with replication across 3 az  
event driven programming with DynamoDB streams

each table has a primary key(partition key, or partition key and sort key)  
each row(item) is made of attributes, max size of item 400KB

throughputs - rcu(read capacity unit), wcu(write capacity unit)  
can use auto scaling of throughput to meet demand or use "burst credit"  
1 wcu - one write per second for an item upto 1KB in size  
1 rcu - item upto 4KB in size, one strongly consistent read per second or two eventually consistent read per second  
eventually consistent read: read after write, unexpected response because of replication  
strongly consistent read, correct data  
if we exceed rcu or wcu then, ProvisionedThroughputExceededExceptions

PutItem - write data (create or replace) to dynamodb  
UpdateItem - update attribute of existing record  
Conditional write - write only if (to deal with concurrency)  
DeleteItem - delete a record, optional condition  
DeleteTable - delete whole table and its items  
BatchWriteItem - 25 put or delete in one call, upto 16 MB of data or 400 KB per item  
GetItem - eventually consistent read by default, ProjectionExpression to retrieve only certain attributes  
BatchGetItem - upto 100 items, 16MB of data

Query: partition key (must be equal), sort key (could be comparison), filter expression (client side filtering)  
upto 1MB of data or number items specified in limit  
can query a table, local secondary index or global secondary index  
efficient

Scan: uses entire table to filter out data, consumes a lot of rcu  
return upto 1MB of data, pagination to continue reading

parallel scans - faster but consumes even more rcu, use limit to mitigate

Local Secondary index - alternate sort key(range key) of a partition, 5 LSI per table, defined at creation

Global Secondary index - secondary partition key + (optional) sort key on the whole table

this new table can project all attributes or main table's primary and sort keys only or some additional attributes and the main table's primary keys

must define rcu and wcu for the gsi, possible to add later or modify (unlike lsi)

optimistic locking/concurrency - conditional write or delete, check version (as opposed to performing a lock)

Dynamodb accelerator - dax, cache for dynamodb, write through

solves hotkey (too many reads) problem, 5 min ttl by default

upto 10 nodes in cluster, multi az, secure

Atomic counters can be implemented using updateItem

DynamoDB streams - change in db, can be read by Lambda or EC2

24 hours of retention, aws managed shards

Kinesis adapter can be used

when the table is modified send:

- KEYS\_ONLY - Only the key attributes of the modified item.

- NEW\_IMAGE - The entire item, as it appears after it was modified.

- OLD\_IMAGE - The entire item, as it appeared before it was modified.

- NEW\_AND\_OLD\_IMAGES - Both the new and the old images of the item.

event source mapping - to read(poll) from stream

dynamodb ttl - delete after expiry date time

to reduce storage and manage table size over time

ttl number - epoch timestamp

transaction - ability to create / update / delete multiple rows in different tables at the same time

all or nothing type of operation, consumes 2x wcu/rcu, transactional read mode

dynamodb as cache for session state (vs elasticache)

write sharding - add a suffix to partition key so that data is distributed

write types - atomic write, batch write, conditional write(in case of concurrent write)

table cleanup - drop table and recreate it

copying table - aws datapipeline(EMR), or backup and restore

RequestLimitExceeded - provisioned throughput exceeds the current throughput limit for your account  
ThrottlingException - rate of requests exceeds the allowed throughput  
ProvisionedThroughputExceededException - rcu or wcu exceeded or the provisioned write capacity for the global secondary index is less than the write capacity of the base table

## API Gateway:

expose rest api

AWS Lambda + API Gateway: No infrastructure to manage, Support for the WebSocket Protocol

Handle API versioning (v1, v2...), different environments (dev, test, prod...), security (Authentication and Authorization)

Create API keys, handle request throttling, Swagger / Open API import to quickly define APIs

Transform and validate requests and responses, Generate SDK and API specifications, Cache API responses

endpoint types:

edge-optimized(default): for global clients, requests routed through cloudfront edge locations, api gateway is in one region

regional: for clients in same region

private: accessed within vpc

deployment stages: api breaking changes, stage variables and lambda alias

linear

canary deployment

all at once

integration type: integrate with backend

Mock - returns a response without sending the request to backend

HTTP - configure integration request and integration response, data mapping using mapping template(vtl), sqs

AWS - configure integration request and integration response, data mapping using mapping template(vtl), for Lambda custom integration

AWS\_Proxy - forward the request json payload, receive response json payload, lambda proxy

HTTP\_Proxy - forward the request/response, application load balancer

api gateway cache - cache per stage, per method cache setting can override, ttl is 300 seconds

if no invalidate cache policy than any client can invalidate the api cache

A client of your API can invalidate an existing cache entry and reload it from the integration endpoint for individual requests. The client must send a request that contains the Cache-Control: max-age=0 header

you can define a [stage variable](#) in a stage configuration, and then set its value as the URL string of an HTTP integration for a method in your REST API. Later, you can reference the URL string using the associated stage variable name from the API setup. This way, you can use the same API

setup with a different endpoint at each stage by resetting the stage variable value to the corresponding URLs. You can also access [stage variables](#) in the mapping templates, or pass configuration parameters to your AWS Lambda or HTTP backend.

A Lambda authorizer is an API Gateway feature that uses a Lambda function to control access to your API. There are two types of Lambda authorizers:

A token-based Lambda authorizer (also called a TOKEN authorizer) receives the caller's identity in a bearer token, such as a JSON Web Token (JWT) or an OAuth token.

A request parameter-based Lambda authorizer (also called a REQUEST authorizer) receives the caller's identity in a combination of headers, query string parameters, stageVariables, and \$context variables.

usage plan:

- who can access one or more deployed API stages and methods

- how much and how fast they can access them

- uses API keys to identify API clients and meter access

- configure throttling limits and quota limits that are enforced on individual client

API keys:

- alphanumeric string values to distribute to your customers, ex -

WBjHxNtoAb4WPKBC7cGm64CBiblb24b4jt8jJHo9

- can use with usage plans to control access

- throttling limits are applied to the API keys

- quotas limits is the overall number of maximum requests

To configure a usage plan

- create one or more APIs, configure the methods to require an API key, and deploy the APIs to stages.

- generate or import API keys to distribute to application developers (your customers) who will be using your API.

- create the usage plan with the desired throttle and quota limits.

- associate API stages and API keys with the usage plan.

cloudwatch logs - info about request/response body

cloudwatch metrics - cache hit, cache miss, count, latency(max 29 seconds then timeout),

integrationLatency, 4xx error (client side), 5xx error (server side)

HTTP Status Code	
400	Bad Request Exception
403	Access Denied Exception, autorisation issue
404	Not Found Exception



409	Conflict Exception
429	Limit Exceeded Exception
429	Too Many Requests Exception
502	Bad Gateway Exception, usually for an incompatible output returned from a Lambda proxy integration backend and occasionally for out-of-order invocations due to heavy loads.
503	Service Unavailable Exception
504	Endpoint Request Timed-out Exception

x-ray - full picture

enable cors to receive api calls from another domain

options pre-flight request must contain headers:

Access-Control-Allow-Methods

Access-Control-Allow-Headers

Access-Control-Allow-Origin

security

iam:

great for users / roles already within your AWS account, + resource policy for cross account

handle authentication + authorization, uses signature v4

custom Authorizer:

great for 3rd party tokens

flexible in terms of what IAM policy is returned

handle Authentication verification + Authorization in the Lambda function

pay per Lambda invocation, results are cached

Cognito User Pool:

manage your own user pool (can be backed by Facebook, Google login etc...)

no need to write any custom code

must implement authorization in the backend

Api gateway integration timeout - 29 sec

Lambda execution timeout - 15 min

## SAM (Serverless Application Model):

configuration using yaml code, creates cloudformation templates

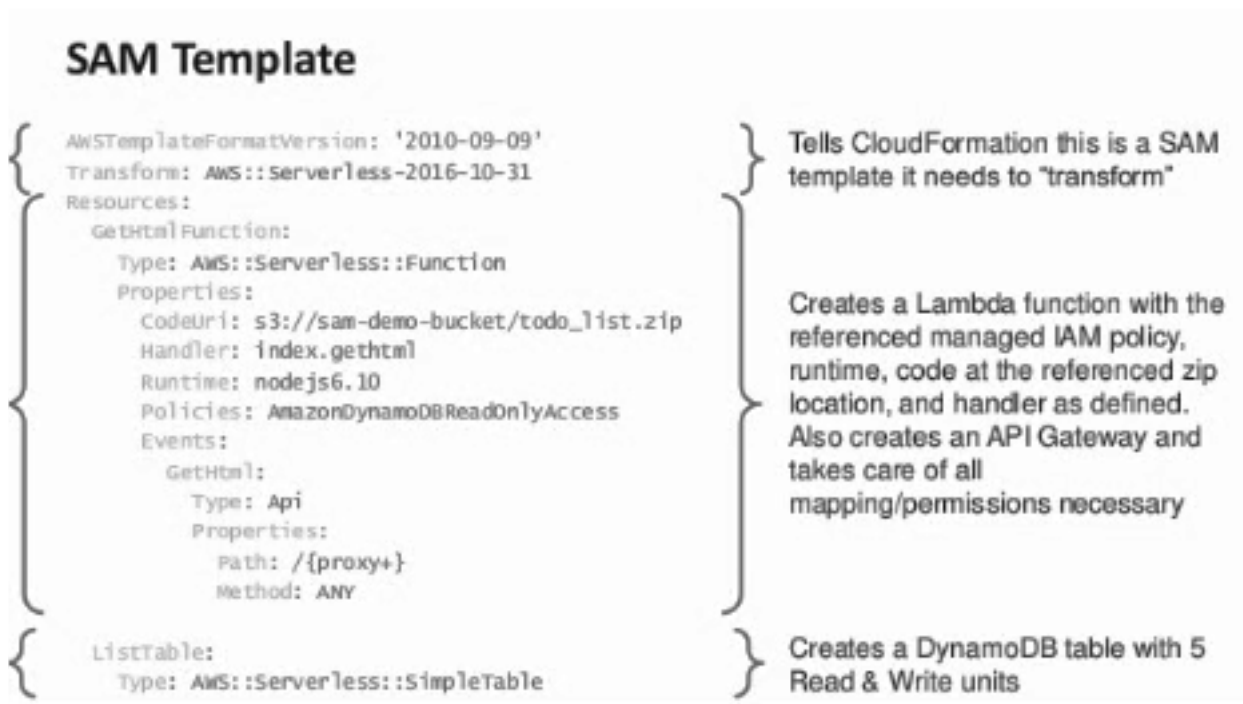
using codedeploy deploy lambda functions

sam can run lambda, api gateway, dynamodb locally

Transform Header indicates it's SAM template, Transform: 'AWS::Serverless-2016-10-31'

Write Code - AWS::Serverless::Function, AWS::Serverless::Api, AWS::Serverless::SimpleTable

Package & Deploy - aws cloudformation package / sam package, aws cloudformation deploy / sam deploy



Node.js and Python functions, you can specify the function code inline in the template, zipfile parameter

Changes to a deployment package in Amazon S3 are not detected automatically during stack updates.

To update the function code, change the object key or version in the template.

## Cognito:

give users an identity, they can interact with our application

cognito user pools: sign in for app users, integrate with api gateway and application load balancer

cognito identity pools (federation identity): users can access aws resources, integrate with cognito user pools as an identity provider

cognito sync (AppSync): synchronize data from device to cognito

cognito user pools - serverless database of mobile and app users, can use federation identity provider  
user pool flow, aws lambda trigger: authentication events, sign up, messages, token creation

cognito identity pools (federation identity) - give users temporary aws credentials  
access private s3 bucket for example

## AppSync:

- combining data from one or more sources
- managed service that uses GraphQL
- retrieve data in real time with WebSocket
- data synchronization

## Step Functions:

- orchestrate workflow, lambda functions, EC2, ECS, on-prem servers, API Gateway
- represent flow as JSON state machine
- sequence, parallel, conditions, timeout, error handling
- max execution time 1 year, manual approval
- retry and catch to handle errors

## AWS STS:

- security token service - grants limited and temporary access to AWS resource (1 hour)
- AssumeRole: Assume roles within your account or cross account
- AssumeRoleWithSAML: return credentials for users logged with SAML
- AssumeRoleWithWebIdentity: return creds for users logged with an IdP (Facebook Login, Google Login, OIDC compatible...), AWS recommends against using this, and using Cognito Identity Pools instead
- GetSessionToken: for MFA, from a user or AWS account root user
- GetFederationToken: obtain temporary creds for a federated user
- GetCallerIdentity: return details about the IAM user or role used in the API call
- DecodeAuthorizationMessage: decode error message when an AWS API is denied

aws:MultiFactorAuthPresent:true

GetSessionToken returns: Access ID, Secret Key, Session Token, Expiration date

If there's an explicit DENY, end decision DENY

If there's an ALLOW, end decision ALLOW Else DENY

IAM policies are attached to users, roles, groups

S3 bucket policies are attached to buckets

when evaluating if an IAM Principal can perform an operation on a bucket, the union of its assigned IAM policies and S3 bucket policies will be evaluated

create dynamic policy with IAM, special policy variable `${aws:username}`

AWS managed policy, Customer managed policy, inline policy (directly in principal)

pass role: iam:passrole

## AWS security and encryption:

encryption in flight - https, ssl enabled website

server side encryption at rest - usually data key is used

client side encryption - envelope encryption/client side data key, server can't decrypt data

KMS - key management service - encryption for an aws service

integrated with IAM for authorization

Amazon EBS: encrypt volumes, Amazon S3: Server side encryption of objects, Amazon Redshift: encryption of data, Amazon RDS: encryption of data, Amazon SSM: Parameter store, etc

symmetric keys (AES-256 keys):

aws services integrated with KMS use this, never get access to the Key unencrypted (must call KMS API to use)

necessary for envelope encryption

asymmetric keys (RSA and ECC key pairs):

public key(encrypt) and private key(decrypt) pair

used for encrypt/decrypt or sign/verify operations

public key downloadable, can't access private key

able to fully manage the keys & policies: Create, Rotation policies, Disable, Enable

able to audit key usage (using CloudTrail)

three types of Customer Master Keys (CMK): AWS Managed Service Default CMK, User Keys created in KMS, User Keys imported (must be 256-bit symmetric key)

for data > 4 KB you can't use KMS, use envelope encryption

To give access to KMS to someone: make sure the Key Policy allows the user, make sure the IAM Policy allows the API calls

kms keys are region bound, copy snapshot across regions - snapshot encrypted with key 1, copy and re encrypt with key 2 in another region, restore in another region which will now use key 2

Default KMS Key Policy:

- created if you don't provide a specific KMS Key Policy

- complete access to the key to the root user = entire AWS account

- gives access to the IAM policies to the KMS key

Custom KMS Key Policy:

- define users, roles that can access the KMS key

define who can administer the key  
useful for cross-account access of your KMS key

Copying Snapshots across accounts:

- create a Snapshot, encrypted with your own CMK
- attach a KMS Key Policy to authorize cross-account access
- share the encrypted snapshot
- (in target) Create a copy of the Snapshot, encrypt it with a KMS Key in your account
- create a volume from the snapshot

encrypt/decrypt api call

if size > 4KB then envelope encryption - GenerateDataKey API, data key dek file, client side encryption using that and then put the encrypted file in envelope along with encrypted data key dek file

AWS Encryption SDK implemented Envelope Encryption for us

Data Key Caching:

- re-use data keys instead of creating new ones for each encryption
- helps with reducing the number of calls to KMS with a security trade-off
- use LocalCryptoMaterialsCache (max age, max bytes, max number of messages)

KMS Request Quotas

when you exceed a request quota, you get a ThrottlingException, to respond, use exponential backoff (backoff and retry)

for cryptographic operations, they share a quota, this includes requests made by AWS on your behalf (ex: SSE-KMS)

for GenerateDataKey, consider using DEK caching from the Encryption SDK

you can request a Request Quotas increase through API or AWS support

can encrypt cloudwatch logs with KMS keys

associate-kms-key : if the log group already exists

create-log-group: if the log group doesn't exist yet

SSM Parameter store: securely store configuration and secrets

can be encrypted using KMS

version tracking, configuration management using path and IAM

notifications with cloudwatch events, integration with cloudformation

SSM Parameter Store Hierarchy

/my-department/

    my-app/

        dev/

            db-url

            db-password

        prod/

db-url  
db-password  
other-app/  
/other-department/  
/aws/reference/secretsmanager/secret\_ID\_in\_Secrets\_Manager  
/aws/service/ami-amazon-linux-latest/amzn2-ami-hvm-x86\_64-gp2

Secrets manager:

capability to force rotation of secrets every X days

automate generation of secrets on rotation (uses Lambda)

integration with Amazon RDS (MySQL, PostgreSQL, Aurora), mostly meant for RDS integration

secrets are encrypted using KMS

Secrets Manager(expensive):

Automatic rotation of secrets with AWS Lambda

Integration with RDS, Redshift, DocumentDB

KMS encryption is mandatory

Can integrate with CloudFormation

SSM Parameter Store:

Simple API

No secret rotation

KMS encryption is optional

Can integrate with CloudFormation

Can pull a Secrets Manager secret using the SSM Parameter Store API

## AWS SES:

send emails using smtp, or aws sdk

can receive emails

## AWS Certificate Manager (ACM):

host public ssl certificates

ACM loads SSL certificates on the following integrations:

Load Balancers (including the ones created by EB)

CloudFront distributions

APIs on API Gateways

Use IAM as a certificate manager only when you must support HTTPS connections in a Region that is not [supported by ACM](#). IAM securely encrypts your private keys and stores the encrypted version in IAM SSL certificate storage. IAM supports deploying server certificates in all Regions, but you must obtain your certificate from an external provider for use with AWS. You cannot

upload an ACM certificate to IAM. Additionally, you cannot manage your certificates from the IAM Console.

## **Amazon Inspector:**

Using Amazon Inspector is an automated security assessment service that helps improve the security and compliance of applications deployed on AWS.

## **Quicksight:**

Amazon QuickSight is a fast, cloud-powered business intelligence service that makes it easy to deliver insights to everyone in your organization. As a fully managed service, QuickSight lets you easily create and publish interactive dashboards that include ML Insights. Dashboards can then be accessed from any device and embedded into your applications, portals, and websites.

## **CloudHSM:**

AWS CloudHSM provides hardware security modules in AWS Cloud. A hardware security module (HSM) is a computing device that processes cryptographic operations and provides secure storage for cryptographic keys.

## **OpsWorks**

It is a configuration management service that provides managed instances of Chef and Puppet.

## **AWS CodeStar**

It provides the tools you need to quickly develop, build, and deploy applications on AWS.

With AWS CodeStar, you can use a variety of project templates to start developing applications on Amazon EC2, AWS Lambda, and AWS Elastic Beanstalk.

The project dashboard in AWS CodeStar makes it easy to centrally monitor application activity and manage day-to-day development tasks such as recent code commits, builds, and deployments.

Because AWS CodeStar integrates with Atlassian JIRA, a third-party issue tracking and project management tool, you can create and manage JIRA issues in the AWS CodeStar dashboard.

### **Whitepapers:**

Serverless architecture with AWS lambda

Running containerised microservices on aws

Microservices on AWS

Api gateway faq

Cloudwatch event faq

